

Discriminant Analysis of Vowel Data

Author: Shaohao Chen , Email: shaohao@bu.edu

I. Description of the problem

Download the vowels data from <https://web.stanford.edu/~hastie/ElemStatLearn/data.html> (<https://web.stanford.edu/~hastie/ElemStatLearn/data.html>). There are 11 vowels to distinguish, based on 10 features for each training/test sample of a vowel. We will use linear and quadratic discriminant methods to analyze the data. The following two ways are used:

1. Quadratic discriminant analysis: Fit a different Gaussian model for each of the 11 classes. This involves determining both a different centroid μ_k and a different covariance Σ_k for each class k .
2. Linear discriminant analysis on a quadratic basis: Increase the input space beyond the 10 features above by including quadratic and interaction terms (e.g. x_i^2 , $x_i x_j$) as well. Then fit the same Gaussian model for the 11 classes. A common covariance Σ is assumed.

II. Implementation with Python

Following is the code implemented by Python. The Numpy package is used for linear algebra operations. The raw vowel data is saved in CSV format and is loaded to Numpy arrays. The QDA and LDA are done by calling functions in the Scikit-learn package. These functions fit the models, predict the results and return correct classification rates. The least-square solver is used for LDA. I wrote a function to construct quadratic basis from the given linear basis. The visualization is done by the Matplotlib. I wrote a function to plot figures for classification results in two-dimensional subspaces. The correct and error predictions in each of the 11 classes are identified.

```

In [2]: # Import necessary python modules.
import numpy as np
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
import matplotlib.pyplot as plt
from matplotlib import colors

# A function to construct quadratic basis from linear basis
def quad_basis(h_lin):
    dim = h_lin.shape[1] # number of original dimensions = number of columns of X
    dim_add = dim*(dim+1)/2 # number of dimensions to be added
    n_row = h_lin.shape[0] # number of rows of X
    h_add = np.zeros((n_row, dim_add)) # shape of the sub matrix for new dimensions
    index=0
    for i in range(0, dim):
        for j in range(i, dim):
            h_add[:,index] = h_lin[:,i]*h_lin[:,j] # compute the sub matrix for new dimensions
            index += 1
    h_quad = np.concatenate((h_lin, h_add), axis=1) # append new dimensions to obtain new X
    return h_quad

# A function to plot results
def plot_classification(y, y_pred, dim1, dim2, p_dim, xlab, ylab, title, n_fig):
    ms=8
    mcolor=['C0', 'C1', 'C2', 'C3', 'C4', 'C5', 'C6', 'C7', 'C8', 'C9']
    correct_i = np.where(y==y_pred) # Index of correct predictions
    error_i = np.where(y!=y_pred) # Index of error predictions
    for i in range(p_dim):
        class_i = np.where(y==i+1) # Index of the i-th class
        correct_class_i = np.intersect1d(correct_i, class_i) # Index of correct predictions in the i-th class
        error_class_i = np.intersect1d(error_i, class_i) # Index of error predictions in the i-th class
        plt.plot(X0[correct_class_i, dim1], X0[correct_class_i, dim2], 'o', markeredgecolor=mcolor[i],
        markerfacecolor='None', markersize=ms) # plot correct predictions
        plt.plot(X0[error_class_i, dim1], X0[error_class_i, dim2], '^', markeredgecolor=mcolor[i], markersize=ms) # plot error predictions
        lsize=18
        plt.xlabel(xlab, fontsize=lsize)
        plt.ylabel(ylab, fontsize=lsize)
        tsize=20
        full_title = "Fig." + n_fig + " " + title
        plt.title(full_title, fontsize=tsize)

# ===== Start main program =====
# Read vowel data from csv file and save in numpy arrays
training_data = np.genfromtxt('vowel_data/training_nohead.csv', delimiter=',')
test_data = np.genfromtxt('vowel_data/test_nohead.csv', delimiter=',')

# Feed X and y with training data
y = training_data[:,1] # 2nd column
X = training_data[:,2:] # from 3rd column onwards
# Feed X0 and y0 with test data
y0 = test_data[:,1] # 2nd column
X0 = test_data[:,2:] # from 3rd column onwards

# Quadratic Discriminant Analysis
qda = QuadraticDiscriminantAnalysis(store_covariance=True) # Call QDA in Scikit-learn
qda.fit(X, y) # Fit with training data
y0_qda = qda.predict(X0) # Predict with test data
print "===== QDA ====="
print "Misclassification rate of training data = ", 1 - qda.score(X, y)
print "Misclassification rate of test data = ", 1 - qda.score(X0, y0)

```

===== QDA =====

Misclassification rate of training data = 0.011363636363636354

Misclassification rate of test data = 0.5281385281385281

===== LDA: quadratic basis =====

Misclassification rate of training data = 0.022727272727272707

Misclassification rate of test data = 0.43939393939393945

Fig.1 QDA results in subspace (x1, x2)

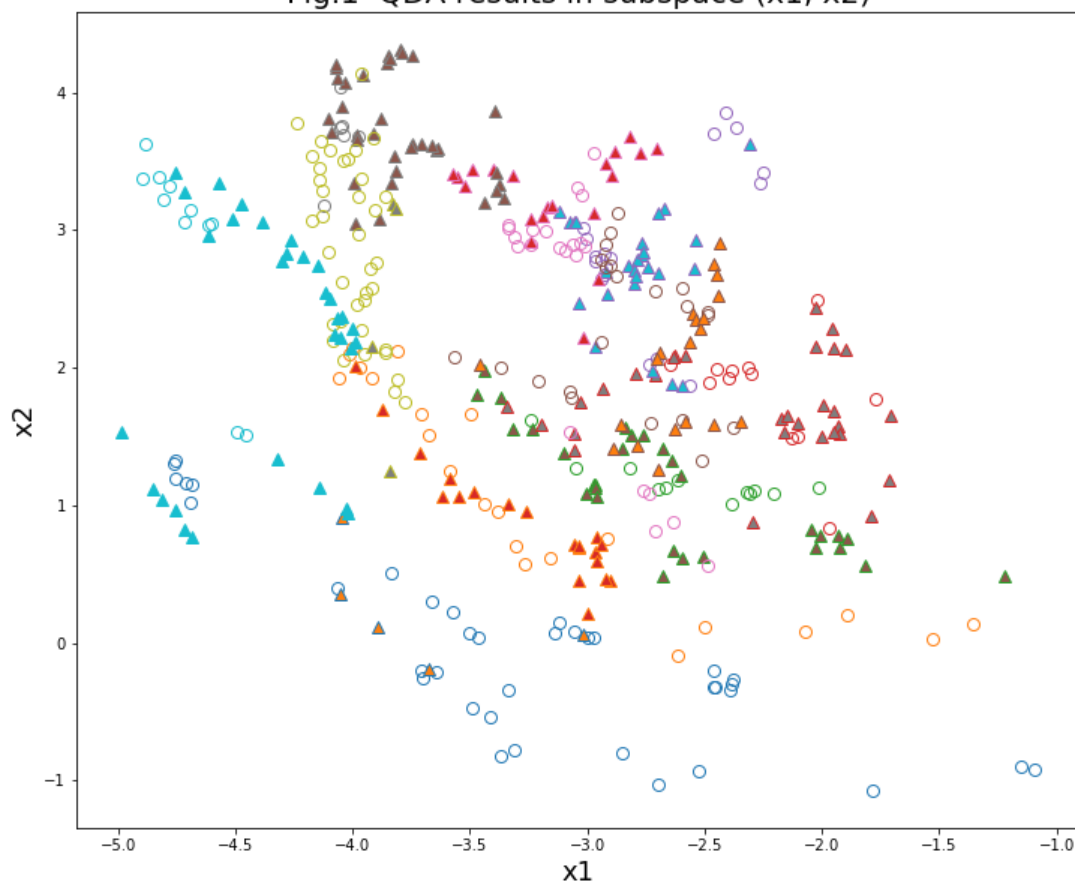


Fig.2 QDA results in subspace (x5, x6)

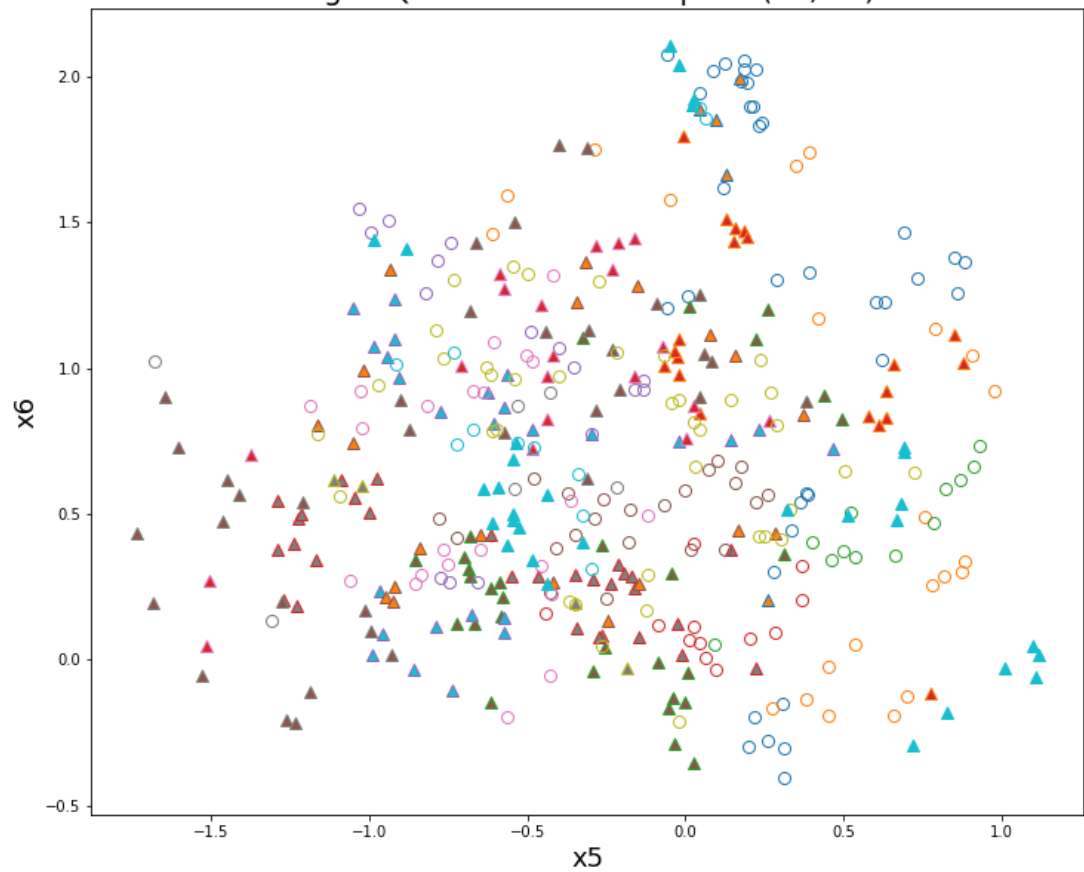


Fig.3 QDA results in subspace (x9, x10)

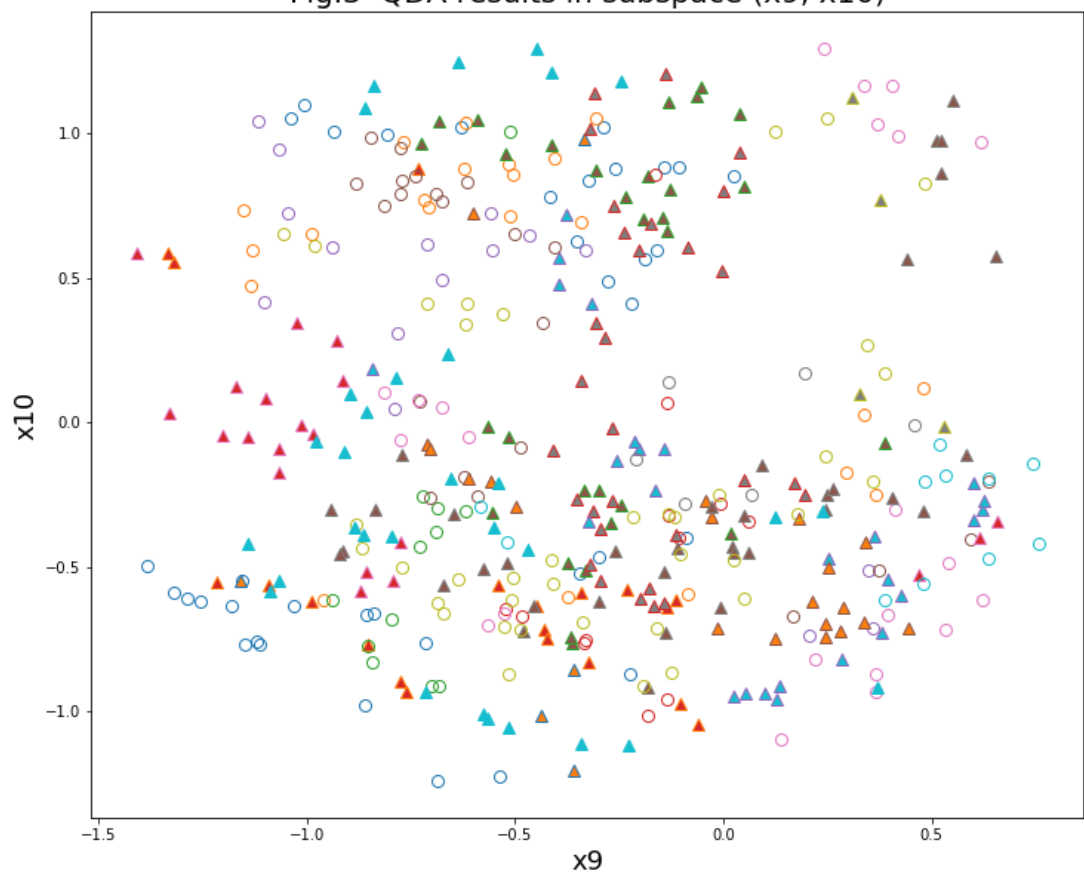


Fig.4 LDA results in subspace (x1, x2)

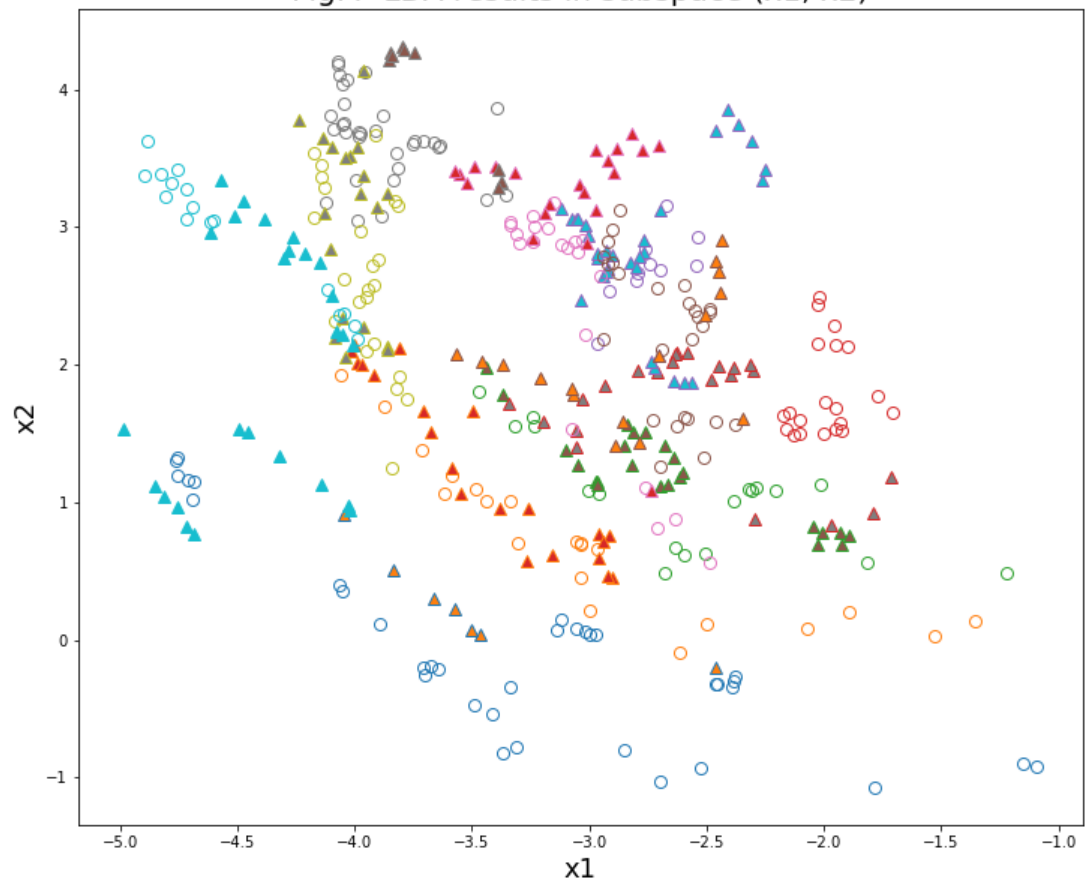


Fig.5 LDA results in subspace (x5, x6)

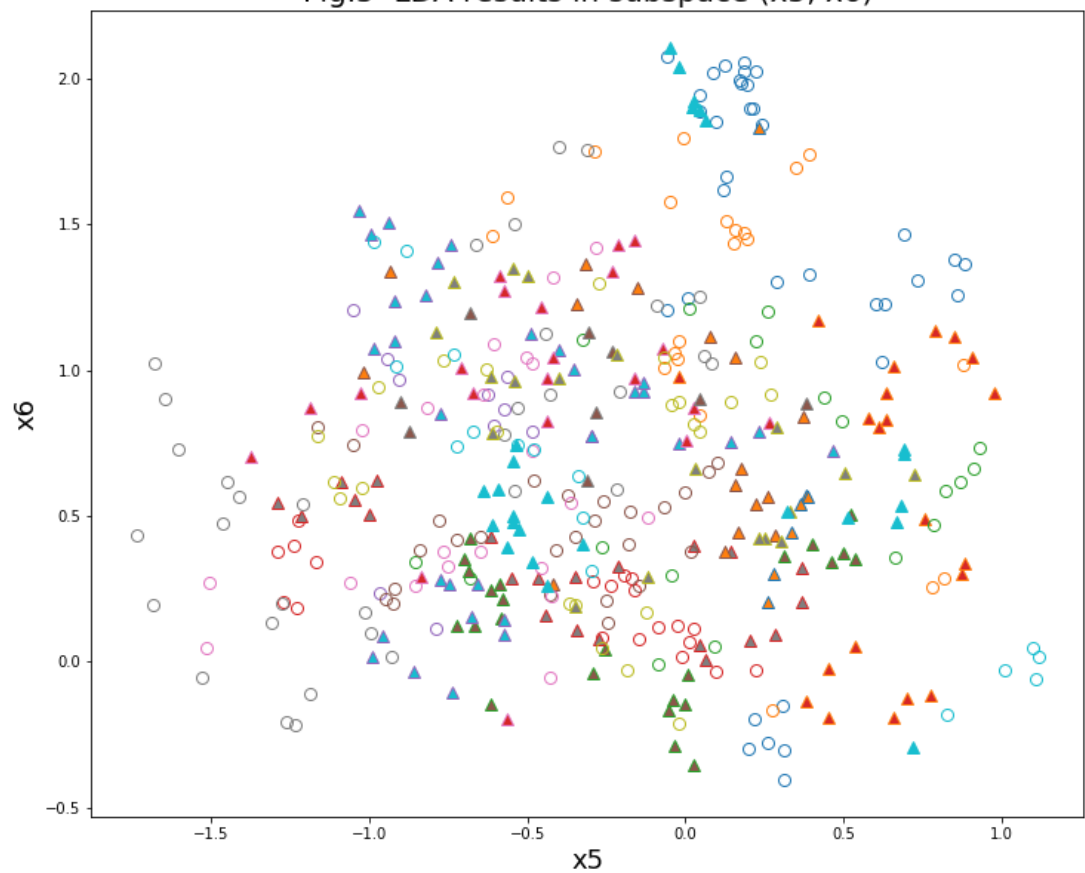
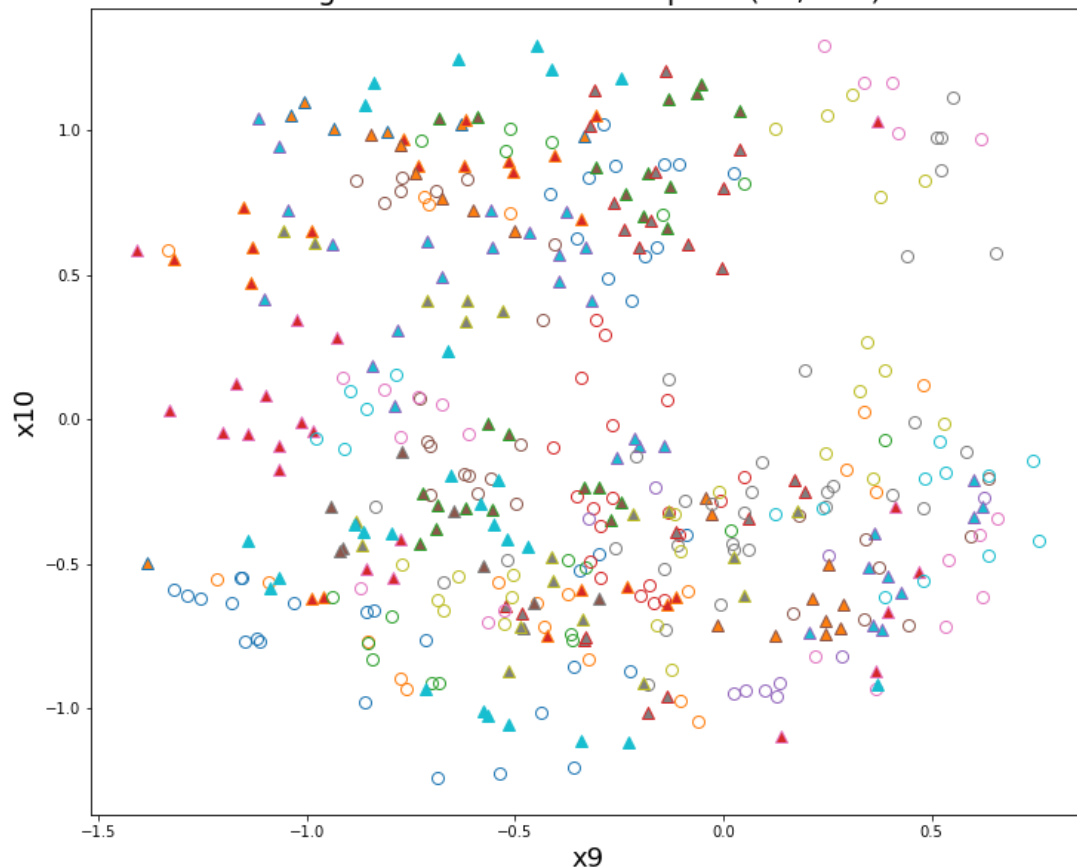


Fig.6 LDA results in subspace (x9, x10)



III. Discussions

III.1 Misclassification rates

For QDA, the misclassification rate of training data is 0.011 and that of test data is 0.528. For LDA on the quadratic basis, the misclassification rate of training data is 0.023 and that of test data is 0.439. As expected, the training misclassification rates are very low, since we use the fitted model to predict with in-sample data. Also as expected, the misclassification rate of the test data is much higher, since we predict with out-of-sample data. The test misclassification rate (0.528) given by the QDA method is the same as the value listed in the Hastie Book and it is acceptable for a 10-dimension 11-class problem. The test misclassification rate given by the LDA method on the quadratic basis (0.439) is better than that of QDA, and is better than the results of many other methods listed in the original document in the vowel data package. In this LDA model, total 65 basis functions are used. To obtain an even better test misclassification rate, cross-validation may be needed.

III.2 Visualization

The data is in 10 dimensions. In both models, the discriminant borders are 10-dimension quadratic hyperplanes. It is not straightforward to visualize the data. I visualize the test data in three selected two-dimension subspaces (See Fig. 1 - Fig. 6 above). Points in the 11 classes are plotted in different colors, while correct and error predicted points are in different shapes of markers. The circle marker is for correct predictions, while the triangle marker is for error predictions.